

Computational Methods Assignment

Task 1

I decided to write a program to brute-force this, rather than do it by hand. Here is the pseudocode:

```
Let adjacency_matrix = {{0,10,15,12,20}, {10,0,12,25,14}, {15,12,0,16,28}, {12,25,16,0,17}, {20,14,28,17,0}}
Let cargo_pickup_weights = [20,40,70,10,30]
```

```
Function get_distance(a, b):
    Return adjacency_matrix[a][b]
End Function
```

```
Function calculate_fuel_cost(a, b, c, d, e):
    Let total_fuel = 0
    Let total_weight = 0

    total_weight = total_weight + cargo_pickup_weights[a]
    total_fuel = total_fuel + get_distance(a,b)*total_weight

    total_weight = total_weight + cargo_pickup_weights[b]
    total_fuel = total_fuel + get_distance(b,c)*total_weight

    total_weight = total_weight + cargo_pickup_weights[c]
    total_fuel = total_fuel + get_distance(c,d)*total_weight

    total_weight = total_weight + cargo_pickup_weights[d]
    total_fuel = total_fuel + get_distance(d,e)*total_weight

    total_weight = total_weight + cargo_pickup_weights[e]

    Return total_fuel*25
End Function
```

```
Let all_possible_sequences = []
```

```
Let a = 0
While a < 5:
```

```

Let sequence = [a]

Let b = 0
While b < 5:
    If sequence Contains b:
        Continue
    End If
    Append b to sequence

    Let c = 0
    While c < 5:
        If sequence Contains c:
            Continue
        End If
        Append c to sequence

        Let d = 0
        While d < 5:
            If sequence Contains d:
                Continue
            End If
            Append d to sequence

            Let e = 0
            While e < 5:
                If sequence Contains e:
                    Continue
                End If
                Append e to sequence

                Append sequence to all_possible_sequences
                sequence = [a,b,c,d]
                Increment e
            End While
            sequence = [a,b,c]
            Increment d
        End While
        sequence = [a,b]
        Increment c
    End While
    sequence = [a]
    Increment b
End While

Increment a
End While

```

```

Open "brute_force.csv" as file
Let index = 0
While index < Length of all_possible_sequences:
    Let seq = all_possible_sequences[index]
    Output seq[0] to file
    Output "," to file
    Output seq[1] to file
    Output "," to file
    Output seq[2] to file
    Output "," to file
    Output seq[3] to file
    Output "," to file
    Output seq[4] to file
    Output "," to file
    Output calculate_fuel_cost(seq[0], seq[1], seq[2], seq[3], seq[4]) to file
    Output "\n" to file
End While

```

```

Close file
Copy

```

And the equivalent python code:

```

alpha = 0
beta = 1
gamma = 2
delta = 3
epsilon = 4

```

```

adjacency_matrix = [[0,10,15,12,20], [10,0,12,25,14], [15,12,0,16,28], [12,25,16,0,17], [20,14,28,17,0]]

```

```

cargo_pickup_weights = [20,40,70,10,30]

```

```

def get_distance(a, b):
    return adjacency_matrix[a][b]

def calculate_fuel_cost(a, b, c, d, e):
    total_fuel = 0
    total_weight = 0

    total_weight += cargo_pickup_weights[a]
    total_fuel += get_distance(a,b)*total_weight

    total_weight += cargo_pickup_weights[b]
    total_fuel += get_distance(b,c)*total_weight

```

```

total_weight += cargo_pickup_weights[c]
total_fuel += get_distance(c,d)*total_weight

total_weight += cargo_pickup_weights[d]
total_fuel += get_distance(d,e)*total_weight

total_weight += cargo_pickup_weights[e]

return total_fuel*25

all_possible_sequences = []

for a in range(5):
    sequence = [a]
    for b in range(5):
        if (b in sequence): continue
        sequence.append(b)
        for c in range(5):
            if (c in sequence): continue
            sequence.append(c)
            for d in range(5):
                if (d in sequence): continue
                sequence.append(d)
                for e in range(5):
                    if (e in sequence): continue
                    sequence.append(e)
                    all_possible_sequences.append(sequence)
                    sequence = [sequence[0], sequence[1], sequence[2], sequence[3]]
                    sequence = [sequence[0], sequence[1], sequence[2]]
                    sequence = [sequence[0], sequence[1]]
                    sequence = [sequence[0]]

csv_data = ""
for seq in all_possible_sequences:
    csv_data += str(seq[0]) + ","
    csv_data += str(seq[1]) + ","
    csv_data += str(seq[2]) + ","
    csv_data += str(seq[3]) + ","
    csv_data += str(seq[4]) + ","
    csv_data += str(calculate_fuel_cost(seq[0], seq[1], seq[2], seq[3], seq[4])) + "\n"

print("generated " + str(len(all_possible_sequences)) + " sequences")

file = open("brute_force.csv", "w")

```

```

file.write(csv_data)
file.close()
Copy

```

Here is the CSV file which is produced by the python program, and a more formatted Excel conversion:

```

brute_force.csv
brute_force.xlsx

```

Reading from these we can see that the cheapest route is:
3 0 4 1 2 = Delta -> Alpha -> Epsilon -> Beta -> Gamma
Costs 69000 intergalactic currency

This approach isn't a good way to find the shortest path since it requires checking the cost of an enormous and rapidly increasing search space. Specifically possible routes, being the number of planets. Factorial time, is a very very bad time complexity and building a route with many destinations would take a very long time. In order to evaluate the cost of each route, we also have to traverse the whole list of planets representing the route, which is length n , so the real time complexity is $O(n!n)$.

Task 2

0	1	2	3	4	5	6	7	8	9	start	end	i	j	pivot value	notes
10	15	12	12	25	16	20	14	28	17	0	9	-1	10	25	partition the whole list
												0	10		
												1	10		
												2	10		
												3	10		
												4	10		
												4	9		
10	15	12	12	17	16	20	14	28	25						swap 4 and 9
												5	9		
												6	9		
												7	9		
												8	9		
												8	8		
												8	7		partitioning finished
10	15	12	12	17	16	20	14			0	7	-1	8	12	subsort first half
												0	8		
												1	8		
												1	7		
												1	6		
												1	5		
												1	4		
												1	3		
10	12	12	15	17	16	20	14								swap 1 and 3
												2	3		
												2	2		
												2	1		partitioning finished

0	1	2	3	4	5	6	7	8	9	start	end	i	j	pivot value	notes
10	12									0	1	-1	2	10	subsort first half of first half
												0	2		
												0	1		
												0	0		subsort done
		12	15	17	16	20	14			2	7	1	8	17	subsort second half of first half
												2	8		
												3	8		
												4	8		
												4	7		
		12	15	14	16	20	17								swap 4 and 7
												5	7		
												6	7		
												6	6		
												6	5		partitioning finished
		12	15	14	16					2	5	1	6	15	subsort first half of second half of
												2	6		
												3	6		
												3	5		
												3	4		
		12	14	15	16										swap 3 and 4
												4	4		
												4	3		partitioning finished
		12	14							2	3	1	4	12	subsort first half of first half of se
												2	4		
												2	3		
												2	2		subsort done
				15	16					4	5	3	6	15	subsort second half of first half of
												4	6		
												4	5		
												4	4		subsort done
						20	17			6	7	5	8	20	subsort second half of second half
												6	8		
												6	7		
						17	20								swap 6 and 7
												7	7		
												7	6		subsort done
								28	25	8	9	7	10	28	subsort second half
												8	10		
												8	9		
															swap 8 and 9
								25	28						
												9	9		
												9	8		subsort done
10	12	12	14	15	16	17	20	25	28						sort done

This trace table represents a quicksort, and below is the pseudocode for it.

Procedure swap(array, first, second) Being:

```
    Let temp = array[first]
    array[first] = array[second]
    array[second] = temp
```

End Procedure

Function partition_array(array, start, end) Begin:

```
    // Place the pivot in the middle, this tends to have better performance
    Let pivot_index = ((end - start)/2 Rounded Down) + start
    Let pivot = array[pivot_index]
```

```
    // Initialise pointers
    Let i = start - 1
    Let j = end + 1
```

While True:

```
    // Increment i then break if the targeted element is swappable
```

While True:

```
    Increment i
    If array[i] >= pivot:
        Break
    End If
```

End While

```
    // Decrement j then break if the targeted element is swappable
```

While True:

```
    Decrement j
    If array[j] <= pivot:
        Break
    End If
```

End While

```
    // Return the partition point if i and j meet/cross, otherwise swap their values
```

If i >= j:

```
    Return j
```

Else:

```
    swap(array, i, j)
```

End If

End While

End Function

Procedure quick_sort(array, start, end) Begin:

```
    // Return if there is nothing to sort
```

If start >= end:

```
    Return
```

```

End If

// Perform first sorting pass over current whole array
Let split_index = partition_array(array, start, end)

// Perform subsorts on partitioned arrays
quick_sort(array, start, split_index)
quick_sort(array, split_index + 1, end)
End Procedure
Copy

```

This is an implementation of the quicksort algorithm, using Hoare's pivot choice and pair-of-pointers method. It makes use of recursive quicksort calls to sort a list by swapping items so that they effectively end up grouped (in each sublist) in groups of larger and smaller items; these sublists can then be sorted using the same method, until there is only one item in each sublist (this is the base case for the recursion). This is an example of a divide-and-conquer approach, as the subsequent quicksorts can be parallelised, since they are independent from one another. Quicksort, depending on implementation (particularly choice of pivot) as well as how sorted data already is, usually has worst-case complexity. However with Hoare's partitioning scheme using the middle-pivot (as opposed to pivoting at the start or end value) tends to have worst-case complexity of .

Task 3

Bearing in mind a greedy strategy chooses the best option in the short term and does not look ahead, I experimented with two different techniques: first, to traverse the graph choosing to move along the **cheapest weighted edge** (excluding any which lead to already visited nodes) at every node; second, to traverse along the edge to the **lowest cargo mass** (using mass here to be distinct from edge weights) **adjacent unvisited** node.

The second of these produced a resulting route (starting at delta, since it has the lowest cargo mass to collect) of **delta -> alpha -> epsilon -> beta -> gamma**, costing **69000** intergalactic Vbucks, which happens to also be the optimal path found by the brute-force method.

The first approach by contrast, starting at the same place, ended up choosing a **delta -> gamma -> beta -> alpha -> epsilon** route, which cost almost double the other method at **126500** intergalactic Vbucks.

It makes sense that a mass-focused route is better, since mass accumulates during the graph traversal, whereas the edge weightings (distances between planets) do not accumulate in the same way.

I originally wrote an implementation which used a bubble sort to pre-sort the planet cargo masses, followed by lots of lookups of planet data, as I was keeping the graph data (edge weights, planet cargo masses, etc) in a set of lists. This resulted in a rather unclear algorithm. I decided to come back to this exercise during a seminar, since I wasn't happy with complexity, and the general unclearness of the algorithm. I rewrote the greedy strategy using the same basic algorithm, but with a better implementation. now planets are stored as **structures**, containing all the information about them and how they connect, meaning the amount of look-ups in lists is significantly reduced.

I found that, surprisingly, when constructing the list of connections that one node (i.e. planet) has with other nodes, it's actually *better* to not bother sorting the list by cargo mass (i.e. to reduce

searching later). This is because the later code not only needs to find the next lowest cargo mass planet, it needs to *find one which has not already been visited*, meaning we end up doing a linear search through the connected planets regardless. This means choosing between either an sorting pass with an traversal pass, or an preparation pass with an traversal pass, the latter of which is clearly better. The pseudocode and C++ implementation for the improved method are below.

```

Let NUM_PLANETS = 5
Let FUEL_COST = 25

// structure holding data about a planet (i.e. a node)
Structure planet
    Let name = ""
    Let index = 0
    Let cargo_mass = 0
    Let links = {}
End Structure

// starting data about the planets
Let node_names = { "alpha", "beta", "gamma", "delta", "epsilon" }
Let cargo_masses = { 20,40,70,10,30 }
Let adjacency_matrix = { {0,10,15,12,20}, {10,0,12,25,14}, {15,12,0,16,28}, {12,25,16,0,17}, {20,14,28,17,0} }

// initialise the nodes in the graph
Let planets = {}
Let i = 0
While i < NUM_PLANETS
    Let p = Create planet
    name Of p = node_names[i]
    cargo_mass Of p = cargo_masses[i]
    index Of p = i
    Append p To planets

    Increment i
End While

Let p_minimum = planets[0]

// setup links between nodes
For p_origin In planets

    // update the lowest cargo mass planet
    // since we want to start at this planet
    // and this saves using a second loop
    If cargo_mass Of p_origin < cargo_mass Of p_minimum
        p_minimum = p_origin
    End If
End For

```

```

    // add links from this node to all other nodes
    // but not itself
    Let i = 0
    While i < NUM_PLANETS
        If planets[i] Not p_origin
            Let distance = adjacency_matrix[i][index Of p_origin]
            Append {planets[i], distance} To links Of p_origin
        End If
        Increment i
    End While
End For

// traverse the graph, keeping track of which planets
// have been visited, and which haven't
Let visited = {}
Fill visited With False NUM_PLANETS Times

Let spaceship_mass = 0
Let fuel_cost = 0
Let sequence = ""

Let p_current = p_minimum
Let p_next Be Empty

Loop Forever

    // find the unvisited planet from the current
    // with the lowest cargo mass, via linear search
    Let minimum_mass = Infinity
    Let d_min_mass = -1
    Let p_min_mass Be Empty

    For l_candidate In links Of p_current
        Let p_candidate = l_candidate[0]
        If visited[index Of p_candidate] = False
            If cargo_mass Of p_candidate < minimum_mass
                minimum_mass = cargo_mass Of p_candidate
                p_min_mass = p_candidate
                d_min_mass = l_candidate[1]
            End If
        End If
    End For

    // if there were no unvisited planets
    // other than the current one, break from the loop
    If p_min_mass Is Empty

```

```

        Break Loop
    End If

    // update the next planet we want to visit
    // this will be the one with the next lowest cargo mass
    p_next = p_min_mass
    d_next = d_min_mass

    // add on the calculate fuel cost for the journey between
    // the current planet and the next
    spaceship_mass = spaceship_mass + cargo_mass Of p_current
    fuel_cost = fuel_cost + spaceship_mass * d_next * FUEL_COST

    // update the sequence string
    sequence = sequence + name Of p_current + " -> "

    // mark this planet as visited
    visited[index Of p_current] = true

    // move onto the next planet
    p_current = p_next
End Loop

// finalise and output the result
sequence = sequence + name Of p_current

Output "Found sequence: " + sequence
Output "Costing: " + fuel_cost
Copy

#include <string>
#include <vector>
#include <iostream>

#define NUM_PLANETS 5
#define FUEL_COST 25

using namespace std;

// data about a planet
struct planet
{
    string name;
    int index;
    int cargo_mass;
    vector<pair<planet*, int>> links;
};

```

```

int main()
{
    // starting data about planets
    string node_names[NUM_PLANETS] = { "alpha", "beta", "gamma", "delta", "epsilon" };
    int cargo_masses[NUM_PLANETS] = { 20,40,70,10,30 };
    int adjacency_matrix[NUM_PLANETS][NUM_PLANETS] = { {0,10,15,12,20}, {10,0,12,25,14}, {15,12,0,16,
    // create nodes
    vector<planet*> planets;
    for (int i = 0; i < NUM_PLANETS; i++)
    {
        planet* p = new planet();
        p->name = node_names[i];
        p->cargo_mass = cargo_masses[i];
        p->index = i;

        planets.push_back(p);
    }

    planet* p_minimum = planets[0];

    // setup links between nodes
    for (planet* p_origin : planets)
    {
        // update lowest cargo planet while we're here
        // saves having another loop
        if (p_origin->cargo_mass < p_minimum->cargo_mass)
        {
            p_minimum = p_origin;
        }

        // add links to other nodes
        // not sorted, since sorting them would
        // actually take more time (O(n^2) inside an n-loop)
        for (int i = 0; i < NUM_PLANETS; i++)
        {
            if (planets[i] == p_origin) continue;

            p_origin->links.push_back(pair<planet*, int>(planets[i], adjacency_matrix[i][p_origin->index]));
        }
    }

    // traverse, keep list of visited
    bool visited[NUM_PLANETS] = { false };

```

```

int spaceship_mass = 0;
int fuel_cost = 0;
string sequence = "";

planet* p_current = p_minimum;
planet* p_next = NULL;
int d_next = 0;

while (true)
{
    // find the unvisited planet with the lowest cargo mass
    int minimum_mass = INT_MAX;
    int d_min_mass = -1;
    planet* p_min_mass = NULL;
    for (pair<planet*, int> l_candidate : p_current->links)
    {
        planet* p_candidate = l_candidate.first;
        if (visited[p_candidate->index]) continue;
        if (p_candidate->cargo_mass < minimum_mass)
        {
            minimum_mass = p_candidate->cargo_mass;
            p_min_mass = p_candidate;
            d_min_mass = l_candidate.second;
        }
    }

    // if there were no unvisited planets,
    // other than the current one, break out
    if (p_min_mass == NULL) break;

    // set the planet we intend to visit
    // next (the one with the lowest cargo mass)
    p_next = p_min_mass;
    d_next = d_min_mass;

    // add on the calculated fuel cost
    spaceship_mass += p_current->cargo_mass;
    fuel_cost += spaceship_mass * d_next * FUEL_COST;

    // update the sequence string
    sequence += p_current->name;
    sequence += " -> ";

    // mark it as visited
    visited[p_current->index] = true;
}

```

```

        // move onto the next
        p_current = p_next;
    }

    // output the result
    sequence += p_current->name;

    cout << "Found sequence: " << sequence << endl;
    cout << "Costing: " << fuel_cost << endl;

    return 0;
}
Copy

```

Output:

```

    Found sequence delta -> alpha -> epsilon -> beta -> gamma
    Costing: 69000

```

The output from the second version of the algorithm is, aside from cosmetic modification, exactly the same as from the earlier version. The resulting algorithmic complexity of this solution is , due to the two occurrences of traversing lists of length , times over.

The only further optimisation I think could be made would be using a system of sorted lookup tables for planet data and cargo masses to eliminate the need for searching for the next lowest *unvisited* cargo mass planet, which we just step through sequentially, removing the need for linear searching at each iteration. This might reduce complexity to in the best case if sorted with quicksort.

Task 4

```

alpha.csv
beta.csv
delta.csv
epsilon.csv
gamma.csv

```

I wrote a short C++ program to produce these tables, in other words I not only followed the dynamic programming approach but also implemented the algorithm. The raw exported CSV files are linked above, and then the assembled and formatted Excel spreadsheet is below:

dynamic_programming.xlsx

My code primarily makes use of a **tree structure** which represents the data which is eventually placed in the table, but which is **more compact and easier to traverse**. I made use of a `std::queue` to keep track of the next block of possible sequences to test, and a `std::map` to keep track of the cheapest version of similar routes (used for carrying forward only the better routes). My tree structure makes use of **pointers** to other nodes allocated on the heap. Code is below:

```

#include <map>
#include <string>

```

```

#include <queue>
#include <iostream>
#include <fstream>

// allows for much easier debugging
#define NODE_ZERO 65

using namespace std;

// only supports up to 255 nodes, since each node reference is only a single byte/char
#define NUM_NODES 5

// data describing the network
const int adjacency[NUM_NODES][NUM_NODES] = { { 0, 10, 15, 12, 20 },
                                                { 10, 0, 12, 25, 14 },
                                                { 15, 12, 0, 16, 28 },
                                                { 12, 25, 16, 0, 17 },
                                                { 20, 14, 28, 17, 0 }
};

const int weight[NUM_NODES] = { 20, 40, 70, 10, 30 };
const string names[NUM_NODES] = { "alpha", "beta", "gamma", "delta", "epsilon" };

// struct containing information about a node in the tree
struct cost_tree_node
{
    int cumulative_cost = 0;
    int cumulative_weight = 0;
    string planets_sequence = "";
    unsigned char last_planet = 0;
    cost_tree_node** children = NULL;
    cost_tree_node* parent = NULL;
};

// sort a string sequence alphabetically, but excluding the first and last characters
string sort_sequence(string seq)
{
    if (seq.length() <= 3) return seq;

    string to_sort = seq;

    bool changed = true;
    while (changed)
    {
        changed = false;
        for (int i = 1; i < to_sort.length() - 2; i++)

```

```

    {
        if (to_sort[i] > to_sort[i + 1])
        {
            changed = true;
            unsigned char tmp = to_sort[i];
            to_sort[i] = to_sort[i + 1];
            to_sort[i + 1] = tmp;
        }
    }
}
return to_sort;
}

// output the cost tree as a table to a file
void write_out_table(cost_tree_node* root)
{
    string output = "prefix,";
    for (int i = NODE_ZERO; i < NODE_ZERO + NUM_NODES; i++)
    {
        output += names[i - NODE_ZERO];
        output += ",";
    }
    output += "\n";

    queue<cost_tree_node*> row_queue;
    row_queue.push(root);

    int block = 0;
    while (!row_queue.empty())
    {
        cost_tree_node* row_starter = row_queue.front();
        row_queue.pop();

        if (row_starter->children == NULL) continue;

        if (row_starter->planets_sequence.length() - 1 > block)
        {
            for (int i = 0; i < NUM_NODES + 1; i++)
            {
                output += " ,";
            }
            output += "\n";
            block = row_starter->planets_sequence.length() - 1;
        }

        for (unsigned char c : row_starter->planets_sequence)

```



```

        output += toupper(names[c - NODE_ZERO][0]);
output += ",";

for (int i = 0; i < NUM_NODES; i++)
{
    if (row_starter->children[i] == NULL)
    {
        output += "-,";
        continue;
    }
    output += to_string(row_starter->children[i]->cumulative_cost);
    output += ",";
    row_queue.push(row_starter->children[i]);
}
output += "\n";
}

ofstream file;
file.open(names[root->planets_sequence[0] - NODE_ZERO] + ".csv");
file << output;
file.close();
}

// build the cost tree, this is the actual dynamic programming bit
cost_tree_node* build_dynamic_cost_tree(unsigned char start_node_index)
{
    // make the specified starting node be the root of the tree
    string root_sequence; root_sequence.push_back(start_node_index);
    cost_tree_node* root = new cost_tree_node
    {
        0,
        weight[start_node_index - NODE_ZERO],
        root_sequence,
        start_node_index,
        NULL,
        NULL
    };

    // nodes that need to have their children populated in this block
    queue<cost_tree_node*> this_block_nodes;

    // new child nodes which are the best route starting at string[0] and ending at string[-1]
    // i.e. these are the best (cheapest) permutations of a sequence of planets
    map<string, cost_tree_node*> next_block_routes;

    this_block_nodes.push(root);
}

```

```

// repeat until we reach a block containing cells representing entire routes through the network
for (int block = 0; block < NUM_NODES - 1; block++)
{
    // populate all the rows in the current block
    while (!this_block_nodes.empty())
    {
        // populate the children of a node
        // the parent represents the row label on the left side of a table
        cost_tree_node* parent = this_block_nodes.front();
        this_block_nodes.pop();

        parent->children = new cost_tree_node * [NUM_NODES];

        // calculate the costs of each possible child node (table cell) from the current parent (
        for (unsigned char c = NODE_ZERO; c < NUM_NODES + NODE_ZERO; c++)
        {
            if (parent->planets_sequence.find(c) != string::npos)
            {
                // discard if the sequence has duplicate planets
                parent->children[c - NODE_ZERO] = NULL;
            }
            else
            {
                // create a new child node (table cell) and calculate its cumulative weight and c
                string node_sequence = parent->planets_sequence;
                node_sequence += c;
                cost_tree_node* node = new cost_tree_node
                {
                    parent->cumulative_cost + (parent->cumulative_weight * adjacency[parent->last
                    parent->cumulative_weight + weight[c - NODE_ZERO],
                    node_sequence,
                    c,
                    NULL,
                    parent
                };
                parent->children[c - NODE_ZERO] = node;
                string sorted_seq = sort_sequence(node->planets_sequence);
                if (block >= 2)
                {
                    // check to see if this node represents the cheapest way to travel between
                    // its set of planets, with the same start and end points
                    auto current_best = next_block_routes.find(sorted_seq);
                    // if there are no other routes like this, it must be the best
                    if (current_best == next_block_routes.end()) next_block_routes.insert({ sorted
                    // if there are other routes and this one is the cheapest, update it as the c

```

```

        // so that it gets computed in the next block
        else if (node->cumulative_cost < (*current_best).second->cumulative_cost) next
        // otherwise discard it
    }
    else
    {
        // add the node to the map so that we will compute its children in the next b
        next_block_routes.insert({ sorted_seq, node });
    }
}

}

// queue up the best routes (table cells) from the last block for evaluation in the next one
// where they now become the table rows
for (pair<string, cost_tree_node*> pr : next_block_routes)
{
    this_block_nodes.push(pr.second);
}

// clear and start again
next_block_routes.clear();
}

// write the node tree out as a table to a file
write_out_table(root);

// finally iterate over the list of best routes (table cells) in the last block and find the chea
cost_tree_node* best_route_through_table = this_block_nodes.front();
while (!this_block_nodes.empty())
{
    cost_tree_node* front = this_block_nodes.front();
    this_block_nodes.pop();
    if (front->cumulative_cost < best_route_through_table->cumulative_cost)
    {
        best_route_through_table = front;
    }
}

// return the node describing the best (cheapest) way of traversing the graph, starting at the sp
return best_route_through_table;
}

int main()
{

```

```

for (int i = NODE_ZERO; i < NUM_NODES + NODE_ZERO; i++)
{
    cost_tree_node* res = build_dynamic_cost_tree(i);
    cout << res->cumulative_cost * 25 << endl;
    for (unsigned char c : res->planets_sequence) cout << names[c - NODE_ZERO] << " ";
    cout << endl << endl;
}
}
Copy

```

As can be seen from the console output of the code, by looking at the *lowest cost table cell* in the *last block of each table* (I'm defining a block as a set of rows which have the same number of previously visited planets shown in the far left column, so block 0 has 'A' in the left column, block 1 will have 'AB', 'AG', 'AD', 'AE'), we can find the cheapest route starting at the origin node of the table:

- starting at alpha: 69750 (alpha -> delta -> epsilon -> beta -> gamma)
- starting at beta: 105250 (beta -> epsilon -> delta -> alpha -> gamma)
- starting at gamma: 12600 (gamma -> delta -> alpha -> beta -> epsilon)
- starting at delta: 69000 (delta -> alpha -> epsilon -> beta -> gamma)
- starting at epsilon: 69750 (epsilon -> delta -> alpha -> beta -> gamma)

The *best route overall* can be found by taking the cheapest of these optimal routes, DAEBG for 69000. This is the same optimal route found by brute force, as we would expect (in fact, we can verify that the optimal route costs starting from other planets are also the best routes found starting from those planets by looking at the results of the brute force method).

This dynamic approach is guaranteed to find the optimal route, because we only prune routes which visit the **same planets** (and thus have the same weight), and **end at the same planet** (i.e. have the same options/edge costs for future traversal steps) but with a **worse cost than other routes covering the same planets**.

In terms of complexity, it's evident to see that this is faster than the brute force approach, for two reasons, which correspond to the two main techniques the dynamic approach uses:

1. Memoisation - each time we want to calculate the cost of traversing from one node to another, we don't recalculate the entire cost, just the progressive cost, and previous calculations are saved and reused (reduces time cost to calculate multiple branching routes)
2. Pruning - by pruning provably inferior routes at early stages, we massively reduce the search space. in fact, we reduce our search space all the way down to just 60 full routes covered, from 120 before.

Writing code for this allowed me to test with different numbers of nodes:

n	routes checked to completion	nodes evaluated	total possible routes	nodes evaluated in brute force (equivalent)
5	60	260	120	600
6	120	990	720	4320
7	210	3402	5040	35280
8	336	10808	40320	322560
9	504	32328	362880	3265920

With this table we can see the huge benefit to pruning compared with the brute force approach. The pattern formed is that the number of routes checked to completion is $\frac{1}{2^n}$ when n is the index of the block. This makes sense since at each step, we prune such that the number of routes to examine in the next block is halved, then thirded, etc, leaving only $\frac{1}{2^n}$ routes checked to completion.

We can also see, with some calculation, that the number of actual evaluations (i.e. calculating the cost of a node, and deciding if we should prune it or carry it forward) is equal to $\frac{1}{2^n}$; this represents the total number of rows in the table multiplied by the number of filled cells in each row, block by block (where n is the index of the block). This can be simplified to $\frac{1}{2^n}$, and our pruning becomes even clearer, as we're multiplying the total number of routes by summed fractions, where each fraction is representing $\frac{1}{2^n}$ divided by the ratio of nodes we can prune at each step in the table. the equivalent in brute-force is just the number of routes multiplied by the number of nodes to represent the time taken to calculate the cost of a particular route (n routes, each of length n), so n^2 . Again we can see that our algorithm is much, much better than brute force in terms of complexity.

We need to consider the complexity of the process of checking for alternative routes with the same nodes ('ABGD' vs 'AGBD'). My implementation uses a simple bubble sort, so we can say that overall this implementation has a time complexity of $O(n^2)$. The algorithm could be improved with the use of a better method for detecting permuted sequences of planets (ABGD vs AGBD) which doesn't use sorting but instead hashes the sequence (which could be done in linear $O(n)$ time).

Task 5 - Art Gallery Problem

The art gallery problem is a geometric problem in which an uneven, concave polygon (i.e. 2D shape, though the problem also exists for 3D polyhedra but is much harder to solve) must have the minimum possible number of 'guards' posted at discrete points on or within the polygon such that the entire polygon is 'visible' to the guards (i.e. there is an unbroken ray that leads from any point on any edge to at least one guard). The problem specifically is finding the minimum number of guards.

The analogy is referential to an art gallery of course, with various rooms of different shapes, which is likely to be concave and possible have disconnected obstacles within the gallery. In this scenario, we need to be able to see the whole gallery at once to keep the artworks safe from theft or vandalism, while hopefully minimising the number of guards required to guard it at once. We assume that each guard has 360 degree vision.

Václav Chvátal showed that the maximum possible number of guards required was equal to $\frac{n}{3}$, where n is the number of vertices in the polygon. It can be seen that a single guard must be able to observe the whole of a convex shape (of which a triangle is the simplest and always convex), since no matter where within or on a triangle we place our observation point, we can draw direct lines to all the corners of the triangle. We can see therefore that Chvátal's theorem is true, because if we triangulate the polygon (i.e. we ensure the entire shape is made up of shapes of only 3 vertices) the maximum possible number of triangles we could end up with is $\frac{n-2}{2}$, for the case where the gallery is made up of a number of disconnected triangles which share no vertices with one another, and since we need exactly one guard per triangle, we can then see that the maximum possible number of guards needed would be also $\frac{n}{3}$ [1].

However, our number of guards is reduced when we consider that many triangles will share at least one vertex with a neighbour, usually sharing two with any particular neighbour, which reduces the number of guards required, saving one guard each time this happens. The difficulty of the art gallery problem lies in decomposing a complex shape with many 'ins and outs' and placing guards optimally.

Continuing this train of thought, it is true that we can triangulate our polygon and then colour each of its vertices one of three colours, such that all triangles have one of each of the three colours on its vertices. Steve Fisk points out that by taking the total number of vertices of a certain colour, specifically the colour with the fewest instances in the polygon (i.e. in a polygon with 2 red, 1 green and 1 blue vertices, we would take either 1 green or 1 blue) we can reduce our maximum number of guards required^[2]. This is essentially a mathematical form of the 'sharing vertices' concept described in the previous paragraph.

Both of these geometric proofs are useful for reducing the search space in terms of finding solutions for smaller numbers of guards by setting an upper bound. However, these approaches are somewhat naive as they cannot optimise concave shapes where vertices are not shared. See the diagram below:

Chvatal's proof would tell us that we need a maximum of three guards (since we have seven total vertices, and we have to round up), and Fisk's proof and the colouring scheme tells us that we need at most two guards: these could be placed at the two green vertices, or the two blue ones. However, looking from above, we can clearly see that we only need to place a single guard at the highlighted green vertex. Everything that the other green vertex can see, can also be seen by the highlighted vertex, plus a bit more.

An algorithm to optimise this problem to minimise the number of guards would need to be able to look at different combinations of guard placements to see if the number of guards can be reduced (i.e. brute-force). Heuristics could be applied for example by counting around vertices and looking at their angles relative to the origin vertex to see if there are occluded (invisible from that point) areas. Approximation methods might use a grid to check the coverage of the polygon from certain vertices in the shape.

It's important to note that we have an additional constraint in this problem: keeping guards on vertices. However, there are variations of the problem (and indeed, real-world applications like lighting a stage would be less constrained) which allow guards to be placed on edges, or even anywhere within the polygon, massively (ordinally) increasing the search space, by way of the number of possible configurations.

One approach presented by Ghosh is to reduce the the overall polygon to a set of convex polygons, each of which may be observed by a single guard^[3]. However, even this may not produce optimal results, see the diagram above once again.

The problem, depending on constraints, is considered NP-hard, meaning it's both difficult to solve and difficult to verify in polynomial (i.e.) time.

-
1. Chvátal, V. (2004) '*A combinatorial theorem in plane geometry*', *Journal of Combinatorial Theory, Series B*. Available at: <https://www.sciencedirect.com/science/article/pii/S0095895675900611> (Accessed: 29 October 2023).
 2. Aigner, M., Ziegler, G.M. (2018). '*How to guard a museum. In: Proofs from THE BOOK*'. Springer, Berlin, Heidelberg. https://doi.org/10.1007/978-3-662-57265-8_40
 3. Ghosh, S. K. (1987), '*Approximation algorithms for art gallery problems*', *Proc. Canadian Information Processing Society Congress*, pp. 429–434.